

K-Nearest Neighbors (K-NN) Algorithm: Implementation and Analysis

The implementation includes:

- Distance calculation using Euclidean and Manhattan metrics.
- A fit method that stores the training data.
- A predict method that finds the k k -nearest neighbors and determines the class label.
- A score method to compute classification accuracy.

Data Preparation and Preprocessing:

- The dataset consists of:
 - Training Inputs: X_{train} (features)
 - Training Labels: y_{train} (target class)
 - Test Inputs: X_{test}
 - Test Labels: y_{test}
- Data Normalization:
 - To ensure consistency across feature scales, the StandardScaler from Scikit-learn was used

Results and Analysis

- The model had its lowest accuracy of 72.93% when $k = 1$ and $k = 2$. A smaller k makes the model more sensitive to noise and outliers, which can lead to incorrect classifications.
- As k increases to 3 and 4, the accuracy improves gradually, reaching 77%. This suggests that a slightly larger k helps the model make better decisions by reducing the effect of random noise.
- When k is 5 or higher, the accuracy continues to improve. At $k = 5$, the model reaches 78.88% accuracy, and from this point, accuracy remains stable without large fluctuations. This means the model is finding a good balance between learning from data and avoiding overfitting.
- The highest accuracy of 79.51% is observed when $k = 26$. This suggests that considering more neighbors helps improve predictions. However, if k becomes too large, the model may start losing important details. Large k values also require computing distances for many neighbors, increasing computation time.
- The best value for k is $k = 5$, where the model achieves 78.88% accuracy. This value provides a good balance between accuracy and stability with increasing computational time.